Daniel Hassan

daniel.hassan@example.com | +91 9895552313 | Remote

SUMMARY

NLP and LLM engineer building retrieval-augmented generation systems, semantic search and conversational assistants.

SKILLS

Core: vector DBs, Hugging Face, PyTorch, FastAPI
Tools & Platforms: Git, LangChain, Large Language Models, Retrieval-Augmented Generation
Also familiar with: ML, NLP, Python, Prompt Engineering

EXPERIENCE

Zenith Mobility — Senior LLM Engineer (02/2024 – Present)

- Applied Fine-tuning (LoRA) of open-source Large Language Models, reducing inference cost by 32% versus hosted APIs.
- Designed chunking and embedding strategies for a semantic search platform backed by vector DBs, improving recall@10 by 28%.
- Fine-tuned transformer models with Hugging Face and PyTorch for NLP classification and entity extraction.

Tech: Python, Large Language Models, RAG, LangChain, Fine-tuning

Greenfield Energy Tech — Senior LLM Engineer (08/2019 – 01/2024)

- Applied Fine-tuning (LoRA) of open-source Large Language Models, reducing inference cost by 20% versus hosted APIs.
- Fine-tuned transformer models with Hugging Face and PyTorch for NLP classification and entity extraction.
- Designed chunking and embedding strategies for a semantic search platform backed by vector DBs, improving recall@10 by 34%.

Tech: Python, Large Language Models, Retrieval-Augmented Generation, Git, LangChain

Cobalt Retail Tech — NLP Engineer (11/2017 – 07/2019)

- Fine-tuned transformer models with Hugging Face and PyTorch for NLP classification and entity extraction.
- Applied Fine-tuning (LoRA) of open-source Large Language Models, reducing inference cost by 41% versus hosted APIs.

Tech: Python, Large Language Models, Retrieval-Augmented Generation, LangChain, AWS

ACADEMIC BACKGROUND

- Bachelor of Engineering in Computer Science — VIT Vellore (2013 – 2017)